


More Theory- Based Inference



Kelly McConville

DATA 100

Week 11 | Spring 2026


Announcements

- Exam this week:
 - In-class on Wed
 - Oral on Fri. (Sign-ups are in Slack.)


Goals for Today

- Discuss more random variable-based inference.
- Compare methods of inference.

Statistical Inference Zoom Out – Estimation



Statistical Inference Zoom Out – Testing



Recap:

Central Limit Theorem (CLT): For random samples and a large sample size (n), the sampling distribution of many sample statistics is approximately normal.

- Applies to most of the statistics we have explored so far:
 - Sample mean
 - Sample proportion
 - Sample difference in means
 - Sample difference in proportions
 - Correlation

There Are Several Versions of the CLT!

Response	Explanatory	Numerical_Quantity	Parameter	Statistic
quantitative	-	mean	μ	$\bar{x}$
categorical	-	proportion	p	$\hat{p}$
quantitative	categorical	difference in means	$\mu_1 - \mu_2$	$\bar{x}_1 - \bar{x}_2$
categorical	categorical	difference in proportions	$p_1 - p_2$	$\hat{p}_1 - \hat{p}_2$
quantitative	quantitative	correlation	ρ	r

- Refer to [these tables](#) for:
 - CLT's "large sample" assumption
 - If you want to know the theory-based **R** function.
 - If you are curious about the equations for the test statistic and confidence interval

Let's cover examples of theory-based inference for two variables.

Data Example

We have data on a random sub-sample of the 2010 American Community Survey. The American Community Survey is given every year to a random sample of US residents.

```
1 # Libraries
2 library(tidyverse)
3 library(Lock5Data)
4
5 # Data
6 data(ACS)
7 # Focus on adults
8 ACS_adults <- filter(ACS, Age >= 18)
9
10 glimpse(ACS_adults)
```

Rows: 1,936

Columns: 9

```
$ Sex      <int> 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, ...
$ Age      <int> 38, 18, 21, 55, 51, 28, 46, 80, 62, 41, 37, 42, 69, 48...
$ Married  <int> 1, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 1, 1, 0, 0, ...
$ Income   <dbl> 64.0, 0.0, 4.0, 34.0, 30.0, 13.7, 114.0, 0.0, 0.0, 0.0...
$ HoursWk  <int> 40, 0, 20, 40, 40, 40, 60, 0, 0, 0, 40, 42, 0, 60, 0, ...
$ Race     <fct> white, black, white, other, black, white, white, white...
$ USCitizen <int> 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, ...
$ HealthInsurance <int> 1, 1, 1, 0, 1, 0, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, ...
$ Language <int> 1, 1, 1, 0, 1, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, ...
```


Difference in Proportions

Let's try to determine if there's a relationship between US citizenship and marriage status.

Response variable:

Explanatory variable:


Parameter of interest:

Sample size requirement for theory-based inference:

Difference in Proportions

Let's try to determine if there's a relationship between US citizenship and marriage status.

```
1 # Exploratory data analysis
2 ggplot(data = ACS_adults,
3       mapping = aes(x = factor(USCitizen),
4                     fill = factor(Married))) +
5   geom_bar(position = "fill")
```



```
1 # Sample size
2 ACS_adults %>%
3   count(Married, USCitizen)
```

	Married	USCitizen	n
1	0	0	64
2	0	1	832
3	1	0	79
4	1	1	961

Difference in Proportions

Let's try to determine if there's a relationship between US citizenship and marriage status.

Why is `prop_test()` failing?

```
1 library(infer)
2 ACS_adults %>%
3   prop_test(Married ~ USCitizen,
4             order = c("1", "0"), z = TRUE,
5             success = "1")
```

Error in ``prop_test()``:

! The response variable of ``Married`` is not appropriate since the response variable is expected to be categorical.

Difference in Proportions

Let's try to determine if there's a relationship between US citizenship and marriage status.

```
1 ACS_adults %>%
2   mutate(MarriedCat = case_when(Married == 0 ~ "No",
3                                   Married == 1 ~ "Yes"),
4         USCitizenCat = case_when(USCitizen == 0 ~ "Not citizen",
5                                   USCitizen == 1 ~ "Citizen")) %>%
6   prop_test(MarriedCat ~ USCitizenCat,
7             order = c("Citizen", "Not citizen"), z = TRUE,
8             success = "Yes")
```

```
# A tibble: 1 × 5
  statistic p_value alternative lower_ci upper_ci
  <dbl>    <dbl> <chr>          <dbl>    <dbl>
1   -0.380    0.704 two.sided    -0.101    0.0682
```

Difference in Means

Let's estimate the average hours worked per week between married and unmarried US residents.

Response variable:

Explanatory variable:

Parameter of interest:


Sample size requirement for theory-based inference:

Difference in Means

Let's estimate the average hours worked per week between married and unmarried US residents.

```
1 # Exploratory data analysis
2 ggplot(data = ACS_adults, mapping = aes(x = HoursWk)) +
3   geom_histogram() +
4   facet_wrap(~Married, ncol = 1)
```

```
1 # Sample size
2 ACS_adults %>%
3   drop_na(HoursWk) %>%
4   count(Married)
```



	Married	n
1	0	896
2	1	1040

Difference in Means

Let’s estimate the average hours worked per week between married and unmarried US residents.

Which arguments for `t_test()` reflect my research question?

```
1 library(infer)
2 ACS_adults %>%
3 t_test(HoursWk ~ Married, order = c("1", "0"))
```

```
# A tibble: 1 × 7
  statistic t_df p_value alternative estimate lower_ci upper_ci
  <dbl> <dbl> <dbl> <chr> <dbl> <dbl>
1 4.81 1902. 0.00000160 two.sided 4.55 2.69
6.40
```

```
1 library(infer)
2 ACS_adults %>%
3 t_test(HoursWk ~ Married, order = c("1", "0"),
4       alternative = "greater")
```

```
# A tibble: 1 × 7
  statistic t_df p_value alternative estimate
  <dbl> <dbl> <dbl> <chr> <dbl>
1 4.81 1902. 0.000000800 greater 4.55
2.99 Inf
```

Correlation

We want to determine if age and hours worked per week have a positive linear relationship.

Response variable:

Explanatory variable:


Parameter of interest:

Sample size requirement for theory-based inference:

Correlation

We want to determine if age and hours worked per week have a positive linear relationship.

```
1 # Exploratory data analysis
2 ggplot(data = ACS_adults,
3       mapping = aes(x = Age,
4                     y = HoursWk)) +
5   geom_jitter(alpha = 0.5) +
6   geom_smooth()
```



Correlation

We want to determine if age and hours worked per week have a positive linear relationship.

```
1 cor.test(~ HoursWk + Age, data = ACS_adults, alternative = "greater")
```


Pearson's product-moment correlation

```
data: HoursWk and Age
t = -17.007, df = 1934, p-value = 1
alternative hypothesis: true correlation is greater than 0
95 percent confidence interval:
 -0.3927809  1.0000000
sample estimates:
      cor
-0.360684
```

Correlation

We want to determine if age and hours worked per week have a positive linear relationship.

```
1 # Exploratory data analysis
2 ggplot(data = ACS_adults,
3       mapping = aes(x = Age,
4                     y = HoursWk)) +
5   geom_jitter(alpha = 0.5) +
6   geom_smooth()
```



Have Learned Two Routes to Statistical Inference

Which is **better**?

Is Simulation-Based Inference or Theory-Based Inference better?

Depends on how you define **better**.

- If **better** = Leads to better understanding:
 - Research tends to show students have a better understanding of **p-values** and **confidence** from learning simulation-based methods.
- If **better** = More flexible/robust to assumptions:
 - The simulation-based methods tend to be more flexible but that generally requires learning extensions beyond what we've seen in DATA 100.
- If **better** = More commonly used:
 - Definitely the theory-based methods but the simulation-based methods are becoming more **common**.
- Good to be comfortable with both as you will see both used!

What about inference for linear regression models?

